



# Computer Vision algorithms for R users

useR.2017, Brussels, July 2017

Jan Wijffels

BNOSAC - [jwijffels@bnosac.be](mailto:jwijffels@bnosac.be)

b n o s a c

# Overzicht

- 1 BNOSAC
- 2 Computer Vision - R toolset
- 3 Detect Corners
- 4 Detect Edges
- 5 Detect Lines & Contours
- 6 Identification / track points / predictive image features
- 7 Object detection
- 8 Contact

**b nosac**

open analytical helpers

**BNOSAC**

# [www.bnosc.be](http://www.bnosc.be): R & open source analytics experts

Providing consultancy services in **open source analytical engineering**

- ▶ Support for R / Oracle R Enterprise / Microsoft R / PostgreSQL / Python / ExtJS / Hadoop / ...
- ▶ Expertise in predictive data mining, biostatistics, geostats, R + Python programming, GUI building, artificial intelligence, process automation, analytical web development
- ▶ R implementations, application maintenance & training/consulting
- ▶ Server Pro and Shiny Pro reseller
- ▶ Organise & support RBelgium
- ▶ Hosting CRAN at [www.datatailor.be](http://www.datatailor.be)
- ▶ Contributing to the R community with R packages / R training

## BNOSAC :: OPEN ANALYTICAL HELPERS

BNOSAC provides expertise in statistical modelling, data science, text mining, web scraping, biostatistics, statistical web development and integration services regarding data analytics. It supports all facets of the usage of data analytics at the enterprise. From adhoc analysis to tailor made & integrated solutions. We help set up the best working conditions for data scientists, get them up to speed with training and provide solutions to speed up deployment.



Statistical Services



Data Science



Solutions & Products



In-house training

# R training by BNOSAC: [www.bnoscac.be/training](http://www.bnoscac.be/training)

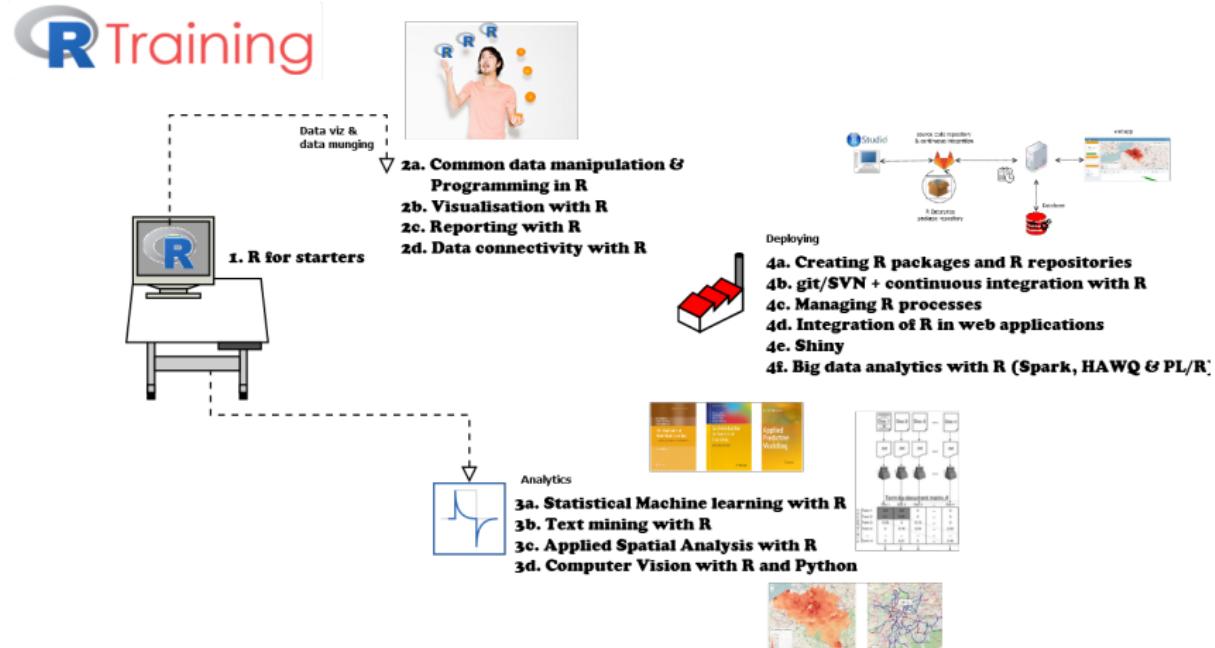


Figure: Training on R / Data Science

**b nosac**

open analytical helpers

# Computer Vision - R toolset

# Computer Vision with R - existing R tools

Quite some packages exist already

## General manipulation and algorithms

**magick** (importing & converting to/from all formats / basic image manipulation); **imager** (interpolation, resizing, warping, filtering, fourier transforms, haar wavelets, morphological operations, denoising, segmentation, gradients, blurring); **EBImage**, used mainly for biological applications; **OpenImageR** (hashing, edge detection, manipulation)

## Domain specific processing

R packages like **adimpro** (smoothing), **radiomics** (texture analysis), **fftw** (fourier transforms), **oro.dicom/oro.nifti** (brain images), **zooimage** (plankton analysis), **wvtool** (wood identification / filters), **Thermimage** (thermal images), **colordistance** (image clustering, color distances), **CRIImage** (tumor detection), **spatstat** (Spatial Point Patterns).

## R is good at interfacing

Rvision/ROpenCVLite (OpenCV from R). API's exist for traditional computer vision (Google Vision API, Microsoft Cognitive Services) or on top of deep learning tools like Keras (**kerasR**), Tensorflow (**tensorflow R package**).

## Typical areas where bnosac.be used image recognition



Figure: Use cases

- ▶ Identify products based on images
- ▶ Quality control of products in factories
- ▶ As input for further processing (predictive models / data enrichment)

## 6 new R packages by bnosac.be

Available at <https://github.com/bnosac/image>

Containing image algorithms lacking in other R packages.

- ▶ **image.CornerDetectionF9**: FAST-9 corner detection (BSD-2).
- ▶ **image.CannyEdges**: Canny Edge Detector (GPL-3).
- ▶ **image.LineSegmentDetector**: Line Segment Detector (LSD) (AGPL-3).
- ▶ **image.ContourDetector**: Unsupervised Smooth Contour Line Detection (AGPL-3).
- ▶ **image.dlib**: Speeded up robust features (SURF) and histogram of oriented gradients (FHOG) features (AGPL-3).
- ▶ **image.darknet**: Image classification using darknet with deep learning models AlexNet, Darknet, VGG-16, GoogleNet and Darknet19. As well object detection using the state-of-the art YOLO detection system (MIT).

More packages and extensions are under development.

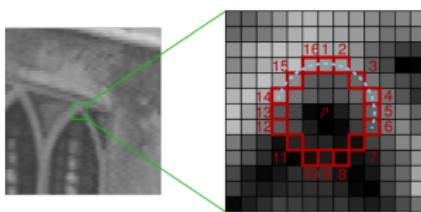
b nosac

open analytical helpers

# Detect Corners

## New R package (1): image.CornerDetectionF9

- ▶ An implementation of the 'FAST-9' **corner detection** algorithm.
- ▶ Finds feature points (corners) in digital images. **These blobs can e.g. be used to track and map objects.**
- ▶ Idea of FAST-9
  - ▶ If a certain number of pixels around that pixel are all brighter or darker than the pixel itself, the pixel is a corner
  - ▶ brighter/darker is defined by a **threshold**



**Figure:** Fast-9 logic - see <https://www.edwardrosten.com/work/fast.html>

- ▶ R package is based on C code at <https://github.com/jcayzac/F9-Corner-Detection-Library>

- ▶ R function `image_detect_corners` requires as input
  - ▶ grayscale matrix with values in 0-255 range + set the threshold
  - ▶ `suppress_non_max`: if pixels are not part of the local maxima they are set to zero (to avoid 2 points in neighbourhood)



Figure: Fast-9 corner detection example

```
library(pixmap)
library(image.CornerDetectionF9)

## Read in a PGM grey scale image
image <- read.pnm(file = system.file("exdata", "chairs.pgm", package="image.CornerDetectionF9"),
                  celres = 1)

## Detect corners
corners <- image_detect_corners(image@grey * 255,
                                 threshold = 100,
                                 suppress_non_max = TRUE)

## Plot the image and the corners
plot(image)
points(corners$x, corners$y, col = "red", pch = 20, lwd = 0.5)
```

b nosac

open analytical helpers

# Detect Edges

## New R package (2): image.CannyEdges

- ▶ Canny edge detector [https://en.wikipedia.org/wiki/Canny\\_edge\\_detector](https://en.wikipedia.org/wiki/Canny_edge_detector),
- ▶ Detects edges in images. Logic:
  - ▶ Remove noise > Find out gradients > Keep only maximum gradient intensities (suppress non-max) > threshold to identify edges > remove edges which are very weak and not connected to strong edges

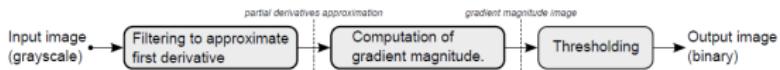


Figure: Edge detection logic - 1st derivative

- ▶ Based on C code from <https://github.com/Neseb/canny>, requiring libpng and fftw3 to be installed (as in `sudo apt-get install libpng-dev fftw3 fftw3-dev pkg-config`)
- ▶ For details on the math: <http://www.ipol.im/pub/art/2015/35>. Other edge detectors in R package OpenImageR.

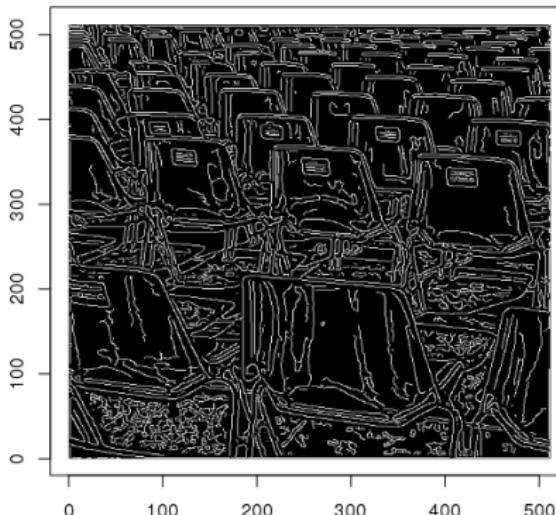
Example below reads in a greyscale image and detects edges in the chairs.

```
library(pixmap)
library(image.CannyEdges)

## Read in a PGM grey scale image
image <- read.pnm(file = system.file("extdata", "chairs.pgm", package="image.CannyEdges"),
                  cellres = 1)

## Detect edges
edges <- image_canny_edge_detector(image@grey * 255, s = 2, low_thr = 3, high_thr = 10)

## Plot the edges
plot(edges)
```



**b nosac**

open analytical helpers

# Detect Lines & Contours

## New R package (3): `image.LineSegmentDetector`

- ▶ **Detection of lines** with the Line Segment Detector (LSD). Idea:
  - ▶ Contour lines are zones of the image where the gray level is changing fast enough from dark to light or the opposite.
  - ▶ Line support regions are being defined by the gradient and grouped into regions with the same direction.
  - ▶ Regions are put in rectangles which define segments

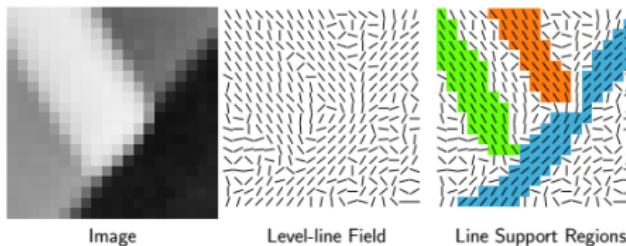


Figure: LSD - line support regions

- ▶ The default arguments provide very good line detections for any image
- ▶ Based on C code available at <https://doi.org/10.5201/ipol.2012.gjmr-lsd>

Example below reads in a greyscale image and detects lines in houses.



Figure: Detect Lines with LSD

```
library(pixmap)
library(image.LineSegmentDetector)

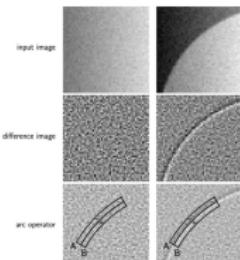
## Read in the PGM file
image <- read.pnm(file = system.file("extdata", "le-piree.pgm", package="image.LineSegmentDetector"),
                   cellres = 1)

## Detect the lines
linesegments <- image_line_segment_detector(image@grey * 255)
linesegments

## Plot the image + add the lines in red
plot(image)
plot(linesegments, add = TRUE, col = "red")
```

## New R package (4): `image.ContourDetector`

- ▶ Detection of contour lines based on paper (Unsupervised Smooth Contour Detection, 2016-11-18, IPOL, <http://www.ipol.im/pub/art/2016/175>). Idea:
  - ▶ remove low frequencies (noise) from the input image.
  - ▶ contours are frontiers separating two adjacent regions, one with significantly larger values than the other.
  - ▶ based on existing edge detector curve candidates are proposed, the curves are piecewise approximated by circular arcs
  - ▶ significance based on non-parametric Mann-Whitney U test to determine whether the samples were drawn from the same distribution or not.
- ▶ The default arguments provide very good contour detections for any image. No need to set detection thresholds like in the Canny edge detection where noise has a big impact on.
- ▶ Based on C code available at <https://doi.org/10.5201/ipol.2016.175>



Example below reads in a greyscale image and detects contour lines in a car.

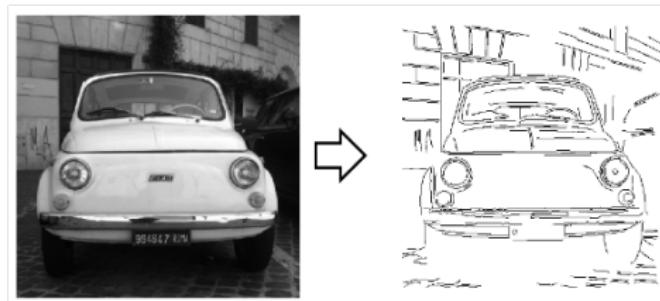


Figure: Detect ContourLines on car

```
library(pixmap)
library(image.ContourDetector)

## Read in the PGM file
image <- read.pnm(file = system.file("extdata", "image.pgm", package="image.ContourDetector"),
                  cellres = 1)

## Detect the contours
contourlines <- image_contour_detector(image@grey * 255)
contourlines

Contour Lines Detector
  found 192 contour lines

## Plot the contour lines
plot(image)
plot(contourlines)
```

Example below reads in a jpg image, converts it to a greyscale image and detects contour lines in the atomium.



**Figure:** Detect ContourLines on the Atomium

```
library(image.ContourDetector)
library(magick)
## Convert jpg to PGM file using the magick package
x <- image_read(path = system.file("extdata", "atomium.jpg", package="image.LineSegmentDetector"))
x <- image_convert(x, format = "pgm", depth = 8)

## Save the PGM file
f <- tempfile(fileext = ".pgm")
image_write(x, path = f, format = "pgm")

## Read in the PGM file, detect the lines
image <- read.pnm(file = f, cellres = 1)
linesegments <- image_line_segment_detector(image@grey * 255)

## Overlay the contour lines on top of the plot
plot(image)
plot(linesegments, add = TRUE, col = "red")
```



open analytical helpers

**Identification /  
track points /  
predictive  
image features**

## New R package (5): image.dlib - SURF

- ▶ Speeded up robust features (**SURF**)
  - ▶ Identifies points in images
  - ▶ Gives a **64-dimensional description of each of the points**
- ▶ SURF descriptors have been used to locate and recognize objects, people or faces, to reconstruct 3D scenes, to track objects and to extract points of interest.
- ▶ The SURF feature descriptor is based on the sum of the Haar wavelet response around the point of interest. Algorithm details at

<http://www.ipol.im/pub/art/2015/69>

- ▶ Want to use it to do object matching/object recognition?
  - ▶ Run SURF descriptor (`image.dlib::image_surf`) on 2 images
  - ▶ Compare the SURF descriptors based on k-nearest-neighbours e.g. with the `rflann` R package (<https://CRAN.R-project.org/package=rflann>).



Figure: SURF for matching

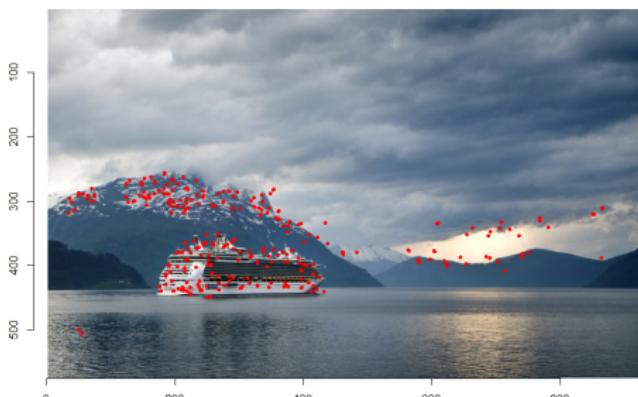
Example below finds the blows. The output shows the points, and the 64-dimensional SURF descriptor for each of the points. If we plot it, we see it finds the boat in the below image and the mountain.

```
library(image.dlib)
f <- system.file("extdata", "cruise_boat.bmp", package="image.dlib")
surf_blobs <- image_surf(f, max_points = 10000, detection_threshold = 50)
str(surf_blobs)

List of 8
 $ points      : num 296
 $ x           : num [1:296] 232 237 282 374 186 ...
 $ y           : num [1:296] 402 371 367 382 416 ...
 $ angle       : num [1:296] -2.99 2.31 2.14 -1.43 1.42 ...
 $ pyramid_scale: num [1:296] 5.27 2.76 2.97 2.94 2.96 ...
 $ score        : num [1:296] 959 630 596 549 526 ...
 $ laplacian    : num [1:296] -1 -1 -1 -1 -1 -1 -1 -1 -1 1 ...
 $ surf         : num [1:296, 1:64] -0.0635 0.1435 0.1229 0.0496 -0.0501 ...

## Plot the points
library(imager)
library(magick)
img <- image_read(path = f)
plot(magick2cimg(img), main = "SURF points")
points(surf_blobs$x, surf_blobs$y, col = "red", pch = 20)
```

SURF points



## New R package (5): image.dlib - FHO

- ▶ Histogram of oriented gradients (**HOG**) features
- ▶ On top of C++ library dlib (<http://dlib.net>) which works with bmp files as input
- ▶ HOG: [http://dlib.net/imaging.html#extract\\_fhog\\_features](http://dlib.net/imaging.html#extract_fhog_features) - popular pedestrian detection algorithm commonly used in the automotive industry
  - ▶ input image is broken into cells that are (cell size) x (cell size) pixels
  - ▶ within each cell we compute a **31 dimensional FHO** vector
  - ▶ this vector describes the gradient structure within the cell and which **can be used in traditional supervised learning**
  - ▶ finds features even in case of changes in illumination and viewpoint, but also due to non-rigid deformations, and intraclass variability in shape and other visual properties.
  - ▶ Object Detection with Discriminatively Trained Part Based Models (P. Felzenszwalb) <http://people.cs.uchicago.edu/~pff/papers/lsvm-pami.pdf>

```
library(image.dlib)
f <- system.file("extdata", "cruise_boat.bmp", package="image.dlib")
x <- image_fhog(f, cell_size = 8)
str(x)

List of 6
 $ hog_height      : num 70
 $ hog_width       : num 116
 $ fhog            : num [1:70, 1:116, 1:31] 0.4 0.4 0.311 0.275 0.399 ...
 $ hog_cell_size   : int 8
 $ filter_rows_padding: int 1
 $ filter_cols_padding: int 1
```

**b nosac**

open analytical helpers

# Object detection

## New R package (6): image.darknet

- ▶ Interface to darknet (<https://pjreddie.com/darknet/yolo>)
- ▶ Applications
  - ▶ Detect locations of objects in an image based on YOLO - You Only Look Once
  - ▶ Classify an image with existing deep learning models (AlexNet, Darknet, VGG-16, Extraction, Darknet19)
- ▶ Darknet: single neural network to the full image.
  - ▶ Network divides the image into regions and predicts bounding boxes and probabilities of objects inside each region.
  - ▶ These bounding boxes are weighted by the predicted probabilities

```
library(image.darknet)

## Define the model
yolo_tiny_voc <- image_darknet_model(
  type = 'detect',
  model = "tiny-yolo-voc.cfg",
  weights = system.file(package="image.darknet", "models", "tiny-yolo-voc.weights"),
  labels = c("aeroplane", "bicycle", "bird", "boat", "bottle", "bus",
            "car", "cat", "chair", "cow", "diningtable", "dog",
            "horse", "motorbike", "person", "pottedplant", "sheep", "sofa", "train", "tvmonitor"))

## Find objects inside the image
image_darknet_detect(file = "imgclairvoyance.jpg", object = yolo_tiny_voc)
```

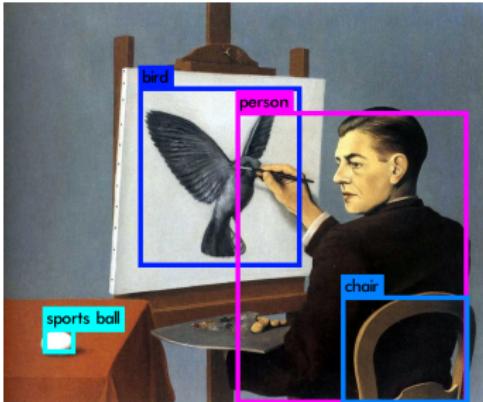


Figure: YOLO - Detection

- ▶ Example was based on YOLO. For other models (AlexNet, Darknet, VGG-16, Extraction, Darknet19), just download the deep learning weights and off you go. Examples at ?image\_darknet\_model

```
weights <- file.path(system.file(package="image.darknet", "models"), "yolo.weights")
download.file(url = "http://pjreddie.com/media/files/yolo.weights", destfile = weights)
```

## ► For classification of an image, use `image_darknet_classify`

```
library(image.darknet)
## Define model
model <- system.file(package="image.darknet", "include", "darknet", "cfg", "tiny.cfg")
weights <- system.file(package="image.darknet", "models", "tiny.weights")
labels <- system.file(package="image.darknet", "include", "darknet", "data", "imagenet.synsetnames.list")
labels <- readLines(labels)

darknet_tiny <- image_darknet_model(type = 'classify',
                                      model = model, weights = weights, labels = labels)

##
## Classify new images alongside the model
##
f <- system.file("include", "darknet", "data", "dog.jpg", package="image.darknet")
x <- image_darknet_classify(file = f, object = darknet_tiny)
x

$file
[1] "C:/Users/Jan/Documents/R/win-library/3.3/image.darknet/include/darknet/data/dog.jpg"

$type
      label probability
1    malamute  0.18206641
2    dogsled   0.12255822
3 Eskimo dog  0.06975935
4     collie   0.04245654
5 Siberian husky  0.03541765
```



Figure: YOLO - for classification

b nosac

open analytical helpers

# Contact

Need more information, send a message at [www.bnoscac.be/index.php/contact/get-in-touch](http://www.bnoscac.be/index.php/contact/get-in-touch)

## OFFICE ADDRESS



The bnoscac office is located in

- La Lustrerie Business Center - <http://www.lalustrerie.be>
- Paleizenstraat 153, 1030 Brussels, Belgium
- close to Gare du Nord in Brussels
- email: info [at] bnoscac [dot] be

**Figure:** Contact [www.bnoscac.be](http://www.bnoscac.be)